

# Employee Turnover Analysis

```
[1]: # Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import (accuracy_score, precision_score, recall_score, f1_score,
                             roc_auc_score, classification_report, confusion_matrix)
from sklearn.cluster import KMeans
from imblearn.over_sampling import SMOTE
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
sns.set_style('whitegrid')
print("Libraries loaded successfully.")

Libraries loaded successfully.
```

## 1. Load Data & Data Quality Checks

Load HR dataset and check for missing values.

```
[4]: # Load dataset (use path relative to notebook location)
df = pd.read_csv('HR_comma_sep.csv')
print("Dataset shape:", df.shape)
print("\nFirst 5 rows:")
df.head()

Dataset shape: (14999, 10)

First 5 rows:
[4]:   satisfaction_level  last_evaluation  number_project  average_montly_hours  time_spend_company  Work_accident  left  promotion_last_5years  sales  salary
      0           0.38          0.53            2                  157                  3             0    1           0    sales    low
      1           0.80          0.86            5                  262                  6             0    1           0    sales  medium
      2           0.11          0.88            7                  272                  4             0    1           0    sales  medium
      3           0.72          0.87            5                  223                  5             0    1           0    sales    low
      4           0.37          0.52            2                  159                  3             0    1           0    sales    low
```

```
[6]: # Data quality checks - Missing values
print("== DATA QUALITY CHECKS ==\n")
print("Missing values per column:")
print(df.isnull().sum())
print("\nTotal missing values:", df.isnull().sum().sum())
print("\nData types:\n", df.dtypes)
print("\nBasic statistics:")
df.describe()

== DATA QUALITY CHECKS ==

Missing values per column:
satisfaction_level      0
last_evaluation          0
number_project           0
average_montly_hours     0
time_spend_company       0
Work_accident            0
left                      0
promotion_last_5years    0
sales                     0
salary                    0
dtype: int64

Total missing values: 0


Data types:
  satisfaction_level      float64
  last_evaluation          float64
  number_project           int64
  average_montly_hours     int64
  time_spend_company       int64
  Work_accident            int64
  left                      int64
  promotion_last_5years    int64
  sales                     object
  salary                   object
dtype: object

Basic statistics:
[6]:   satisfaction_level  last_evaluation  number_project  average_montly_hours  time_spend_company  Work_accident  left  promotion_last_5years
      count      14999.000000  14999.000000  14999.000000  14999.000000  14999.000000  14999.000000  14999.000000  14999.000000
      mean        0.612834    0.716102    3.803054    201.050337    3.498233    0.144610    0.238083    0.021268
      std         0.248631    0.171169    1.232592    49.943099    1.460136    0.351719    0.425924    0.144281
      min         0.090000    0.360000    2.000000    96.000000    2.000000    0.000000    0.000000    0.000000
      25%        0.440000    0.560000    3.000000    156.000000    3.000000    0.000000    0.000000    0.000000
      50%        0.640000    0.720000    4.000000    200.000000    3.000000    0.000000    0.000000    0.000000
      75%        0.820000    0.870000    5.000000    245.000000    4.000000    0.000000    0.000000    0.000000
      max         1.000000    1.000000    7.000000    310.000000    10.000000    1.000000    1.000000    1.000000
```

```
[8]: # Class distribution (target: left)
print("Target 'left' distribution:")
print(df['left'].value_counts())
print("\nPercentage:")
print(df['left'].value_counts(normalize=True).round(3) * 100)
sns.countplot(data=df, x='left')
plt.title('Employee Turnover (0=Stayed, 1=Left)')
plt.show()

Target 'left' distribution:
left
0    11428
1     3571
Name: count, dtype: int64

Percentage:
left
0    76.2
1    23.8
Name: proportion, dtype: float64
```




## 2. Exploratory Data Analysis - Factors Contributing to Turnover


Understand which factors contribute most to employee turnover.

```
[11]: # Correlation with target
df_numeric = df.select_dtypes(include=[np.number])
corr_with_left = df_numeric.corr()['left'].drop('left').sort_values(key=abs, ascending=False)
print("Correlation with 'left' (absolute):\n", corr_with_left)
corr_with_left.plot(kind='barh', figsize=(8,4))
plt.title('Feature correlation with Employee Turnover')
plt.xlabel('Correlation')
plt.tight_layout()
plt.show()


Correlation with 'left' (absolute):
satisfaction_level      -0.388375
Work_accident            -0.154622
time_spend_company       0.144822
average_montly_hours     0.071287
promotion_last_5years   -0.061788
number_project            0.023787
last_evaluation           0.006567
Name: left, dtype: float64
```



```
[13]: # Turnover rate by categorical features
fig, axes = plt.subplots(1, 2, figsize=(12, 4))
df.groupby('salary')['left'].mean().plot(kind='bar', ax=axes[0], color='steelblue')
axes[0].set_title('Turnover rate by Salary')
axes[0].set_ylabel('Turnover rate')
axes[0].tick_params(axis='x', rotation=45)
df.groupby('sales')['left'].mean().plot(kind='bar', ax=axes[1], color='coral')
axes[1].set_title('Turnover rate by Department (sales)')
axes[1].set_ylabel('Turnover rate')
axes[1].tick_params(axis='x', rotation=45)
plt.tight_layout()
plt.show()
```



```
[15]: # Key numeric factors: satisfaction_level, time_spend_company, number_project
fig, axes = plt.subplots(2, 2, figsize=(10, 8))
sns.boxplot(data=df, x='left', y='satisfaction_level', ax=axes[0,0])
axes[0,0].set_title('Satisfaction vs Left')
sns.boxplot(data=df, x='left', y='time_spend_company', ax=axes[0,1])
axes[0,1].set_title('Tenure vs Left')
sns.boxplot(data=df, x='left', y='number_project', ax=axes[1,0])
axes[1,0].set_title('Number of projects vs Left')
sns.boxplot(data=df, x='left', y='average_montly_hours', ax=axes[1,1])
axes[1,1].set_title('Avg monthly hours vs Left')
plt.tight_layout()
plt.show()
print("\nEDA Summary: Low satisfaction, longer tenure, high hours and project count associate with higher turnover.")
```




EDA Summary: Low satisfaction, longer tenure, high hours and project count associate with higher turnover.


### 3. Clustering of Employees Who Left (Satisfaction & Evaluation)

Cluster employees who left based on satisfaction\_level and last\_evaluation.

```
[18]: # Subset: employees who left
df_left = df[df['left'] == 1][['satisfaction_level', 'last_evaluation']].copy()
# Elbow method for optimal k
inertias = []
K = range(2, 8)
for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(df_left)
    inertias.append(kmeans.inertia_)
plt.figure(figsize=(6, 4))
plt.plot(K, inertias, 'bo-')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Inertia')
plt.title('Elbow method for employees who left')
plt.show()
```



```
[20]: # K-Means clustering with k=3 (typical elbow around 3)
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
df_left['cluster'] = kmeans.fit_predict(df_left)
plt.figure(figsize=(8, 5))
sns.scatterplot(data=df_left, x='satisfaction_level', y='last_evaluation', hue='cluster', palette='Set1')
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], s=200, c='black', marker='X', label='Centroids')
plt.title('Clustering of employees who left (Satisfaction vs Evaluation)')
plt.legend()
plt.show()
print("Cluster sizes:", df_left['cluster'].value_counts().sort_index())
```



Cluster sizes: cluster

0	1650
1	977
2	944

Name: count, dtype: int64

### 4. Data Preprocessing & SMOTE for Class Imbalance

Encode categoricals, scale features, and apply SMOTE to balance the 'left' class.

## 4. Data Preprocessing & SMOTE for Class Imbalance

Encode categorical, scale features, and apply SMOTE to balance the 'left' class.

```
[23]: # Encode categorical columns
df_ml = df.copy()
le_sales = LabelEncoder()
le_salary = LabelEncoder()
df_ml['sales_enc'] = le_sales.fit_transform(df_ml['sales'])
df_ml['salary_enc'] = le_salary.fit_transform(df_ml['salary'])
# Features and target
feature_cols = ['satisfaction_level', 'last_evaluation', 'number_project', 'average_montly_hours',
                'time_spend_company', 'Work_accident', 'promotion_last_5years', 'sales_enc', 'salary_enc']
X = df_ml[feature_cols]
y = df_ml['left']
# Train-test split first (SMOTE only on training set)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
print("Before SMOTE - Train:", y_train.value_counts().to_dict())
Before SMOTE - Train: {0: 9142, 1: 2857}

[25]: # Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
# Apply SMOTE on training data only
smote = SMOTE(random_state=42, k_neighbors=5)
X_train_bal, y_train_bal = smote.fit_resample(X_train_scaled, y_train)
print("After SMOTE - Train:", pd.Series(y_train_bal).value_counts().to_dict())
After SMOTE - Train: {0: 9142, 1: 9142}
```

## 5. K-Fold Cross-Validation & Model Training

Train multiple classifiers with stratified k-fold CV and evaluate performance.

```
[28]: # Stratified 5-fold cross-validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42),
    'Random Forest': RandomForestClassifier(n_estimators=100, random_state=42),
    'Gradient Boosting': GradientBoostingClassifier(n_estimators=100, random_state=42)
}
cv_results = {}
for name, model in models.items():
    scores_acc = cross_val_score(model, X_train_bal, y_train_bal, cv=cv, scoring='accuracy')
    scores_f1 = cross_val_score(model, X_train_bal, y_train_bal, cv=cv, scoring='f1')
    scores_roc = cross_val_score(model, X_train_bal, y_train_bal, cv=cv, scoring='roc_auc')
    cv_results[name] = {'accuracy': scores_acc.mean(), 'f1': scores_f1.mean(), 'roc_auc': scores_roc.mean()}
    print(f'{name}: CV Accuracy={scores_acc.mean():.3f}, F1={scores_f1.mean():.3f}, ROC-AUC={scores_roc.mean():.3f}')
cv_df = pd.DataFrame(cv_results).T
cv_df

Logistic Regression: CV Accuracy=0.762, F1=0.768, ROC-AUC=0.817
Random Forest: CV Accuracy=0.990, F1=0.990, ROC-AUC=0.999
Gradient Boosting: CV Accuracy=0.966, F1=0.965, ROC-AUC=0.993
```

	accuracy	f1	roc_auc
Logistic Regression	0.761978	0.767872	0.816927
Random Forest	0.989827	0.989756	0.998775
Gradient Boosting	0.965653	0.965218	0.993267

```
[35]: # Train final models on full balanced train set and evaluate on test set (no SMOTE)
results = []
for name, model in models.items():
    model.fit(X_train_bal, y_train_bal)
    y_pred = model.predict(X_test_scaled)
    results.append({
        'Model': name,
        'Accuracy': accuracy_score(y_test, y_pred),
        'Precision': precision_score(y_test, y_pred, zero_division=0),
        'Recall': recall_score(y_test, y_pred, zero_division=0),
        'F1-Score': f1_score(y_test, y_pred, zero_division=0),
        'ROC-AUC': roc_auc_score(y_test, model.predict_proba(X_test_scaled)[:, 1])
    })
results_df = pd.DataFrame(results)
results_df
```

	Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC
0	Logistic Regression	0.762000	0.500000	0.778711	0.608981	0.822411
1	Random Forest	0.990667	0.994236	0.966387	0.980114	0.992312
2	Gradient Boosting	0.965667	0.921379	0.935574	0.928423	0.989849

## 6. Best Model Selection & Evaluation Metrics Justification

We use **F1-Score** and **ROC-AUC** as primary metrics because turnover is a class-imbalanced problem: we care about correctly identifying employees who will leave (minority class) without excessive false positives. Accuracy alone can be misleading.

```
[31]: # Best model by F1 (and ROC-AUC for robustness)
best_by_f1 = results_df.loc[results_df['F1-Score'].idxmax()]
print("Best model (by F1-Score):", best_by_f1['Model'])
best_model_name = best_by_f1['Model']
best_model = models[best_model_name]
best_model.fit(X_train_bal, y_train_bal)
y_pred_best = best_model.predict(X_test_scaled)
print("\nClassification Report (Best Model):")
print(classification_report(y_test, y_pred_best, target_names=['Stayed', 'Left']))
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_best))
sns.heatmap(confusion_matrix(y_test, y_pred_best), annot=True, fmt='d', cmap='Blues',
            xticklabels=['Stayed', 'Left'], yticklabels=['Stayed', 'Left'])
plt.title(f'Confusion Matrix - {best_model_name}')
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.show()
```

Best model (by F1-Score): Random Forest  
Model Random Forest  
Accuracy 0.990667  
Precision 0.994236  
Recall 0.966387  
F1-Score 0.980114  
ROC-AUC 0.992312  
Name: 1, dtype: object

Classification Report (Best Model):

	precision	recall	f1-score	support
Stayed	0.99	1.00	0.99	2286
Left	0.99	0.97	0.98	714
accuracy			0.99	3000
macro avg	0.99	0.98	0.99	3000
weighted avg	0.99	0.99	0.99	3000

Confusion Matrix:

```
[[2282  4]
 [ 24 690]]
```

Confusion Matrix - Random Forest

## 7. Retention Strategies for Targeted Employees

Based on EDA and clustering, we recommend the following retention strategies for employees at risk of leaving:

```
[33]: # Feature importance (for tree-based best model) to target interventions
if hasattr(best_model, 'feature_importances_'):
    imp = pd.DataFrame({'feature': feature_cols, 'importance': best_model.feature_importances_})
    imp.sort_values('importance', ascending=False)
    imp.plot(x='feature', y='importance', kind='barh', legend=False, figsize=(8, 4))
    plt.title('Feature importance (best model) for retention targeting')
    plt.xlabel('Importance')
    plt.tight_layout()
    plt.show()
    print(imp)
```

Feature importance (best model) for retention targeting

feature	importance
promotion_last_5years	0.293491
time_spend_company	0.254973
number_project	0.153020
average_montly_hours	0.139348
last_evaluation	0.123520
sales_enc	0.020054
salary_enc	0.008201
Work_accident	0.006514
promotion_last_5years	0.000879

## RETENTION STRATEGIES (Targeted by risk factors):

1. LOW SATISFACTION: Conduct stay interviews and pulse surveys; address workload, recognition, and work-life balance.
2. HIGH TENURE + BURNOUT: Offer sabbaticals, role rotation, or project variety; review promotion and growth paths.
3. HIGH HOURS / OVERLOAD: Cap overtime, redistribute projects, and hire to reduce average\_montly\_hours.
4. LOW SALARY (vs market): Benchmark compensation; consider raises or bonuses for high performers.
5. NO PROMOTION (promotion\_last\_5years=0): Clear career ladder and internal mobility programs.
6. DEPARTMENT-SPECIFIC: Tailor interventions by sales/technical/support (e.g., sales: quotas and incentives; technical: learning budget).

Using the model: Score current employees; target retention efforts on those predicted as 'left' with highest probability.